

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import matplotlib.ticker as mtick
%matplotlib inline
```

Loading the data file

```
In [2]: df = pd.read_csv("Customer_churn_data.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleL
0	7590-VHVEG	Female	0	Yes	No	1	No	No p se
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	No p se
4	9237-HQITU	Female	0	No	No	2	Yes	

5 rows × 21 columns



```
In [ ]:
```

```
In [4]: df.columns.values
```

```
Out[4]: array(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',
'TotalCharges', 'Churn'], dtype=object)
```

In [5]: `df.dtypes`

```
Out[5]: customerID      object
gender      object
SeniorCitizen  int64
Partner      object
Dependents    object
tenure      int64
PhoneService  object
MultipleLines  object
InternetService  object
OnlineSecurity  object
OnlineBackup  object
DeviceProtection  object
TechSupport    object
StreamingTV    object
StreamingMovies  object
Contract      object
PaperlessBilling  object
PaymentMethod  object
MonthlyCharges  float64
TotalCharges   object
Churn         object
dtype: object
```

In [6]: `df.describe()`

```
Out[6]:
```

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

```
In [7]: df.isnull().sum()
```

```
Out[7]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport   0
StreamingTV   0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn         0
dtype: int64
```

```
In [8]: df.describe(include = "all")
```

```
Out[8]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
<b>count</b>	7043	7043	7043.000000	7043	7043	7043.000000	7043
<b>unique</b>	7043	2	NaN	2	2	NaN	2
<b>top</b>	7590-VHVEG	Male	NaN	No	No	NaN	Yes
<b>freq</b>	1	3555	NaN	3641	4933	NaN	636
<b>mean</b>	NaN	NaN	0.162147	NaN	NaN	32.371149	NaN
<b>std</b>	NaN	NaN	0.368612	NaN	NaN	24.559481	NaN
<b>min</b>	NaN	NaN	0.000000	NaN	NaN	0.000000	NaN
<b>25%</b>	NaN	NaN	0.000000	NaN	NaN	9.000000	NaN
<b>50%</b>	NaN	NaN	0.000000	NaN	NaN	29.000000	NaN
<b>75%</b>	NaN	NaN	0.000000	NaN	NaN	55.000000	NaN
<b>max</b>	NaN	NaN	1.000000	NaN	NaN	72.000000	NaN

11 rows × 21 columns



```
In [9]: df.describe()
```

Out[9]:

	SeniorCitizen	tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

```
In [10]: df.tail()
```

Out[10]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	
7040	4801-JAZL	Female	0	Yes	Yes	11	No	
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	
7042	3186-AJIEK	Male	0	No	No	66	Yes	

5 rows × 21 columns

```
In [11]: df = df.dropna(how = "all")
```

```
In [12]: df.columns
```

Out[12]: Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'], dtype='object')

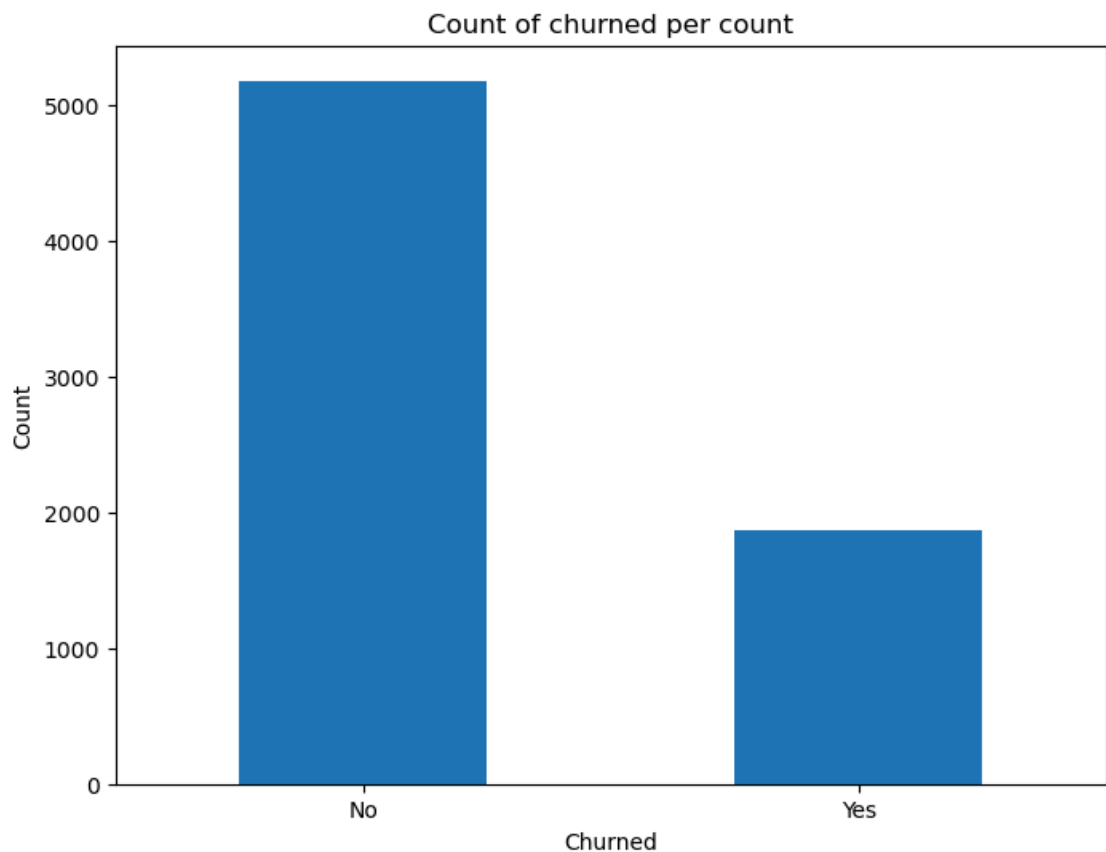
```
In [13]: df["Churn"].value_counts()
```

```
Out[13]: No      5174  
        Yes      1869  
        Name: Churn, dtype: int64
```

```
In [14]: 100*df["Churn"].value_counts()/len(df["Churn"])
```

```
Out[14]: No      73.463013  
        Yes      26.536987  
        Name: Churn, dtype: float64
```

```
In [15]: df["Churn"].value_counts().plot(kind = "bar", figsize = (8,6))  
plt.xlabel("Churned")  
plt.xticks(rotation = "horizontal")  
plt.ylabel("Count")  
plt.title("Count of churned per count")  
plt.show()
```



By above we can see that the data is highly imbalanced in the ratio 73:27 so we analyze the data with other features while taking the largest values separately to get some more insights

In [16]: `df.info(verbose = True)`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

In [17]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [18]: df.isnull().sum()
```

```
Out[18]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport  0
StreamingTV  0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn         0
dtype: int64
```

we don't have missing values

## Data Cleaning

Creating a copy of base data for manipulation and processing

```
In [19]: data = df.copy()
```

Total charges should be numeric amount. Let's convert it into numerical data type

```
In [ ]:
```

```
In [20]: data["TotalCharges"] = pd.to_numeric(data["TotalCharges"], errors = "coerce")
```

```
In [21]: data.isnull().sum()
```

```
Out[21]: customerID      0  
gender      0  
SeniorCitizen  0  
Partner      0  
Dependents    0  
tenure      0  
PhoneService  0  
MultipleLines  0  
InternetService  0  
OnlineSecurity  0  
OnlineBackup  0  
DeviceProtection  0  
TechSupport   0  
StreamingTV   0  
StreamingMovies  0  
Contract      0  
PaperlessBilling  0  
PaymentMethod  0  
MonthlyCharges  0  
TotalCharges   11  
Churn          0  
dtype: int64
```

As we can see there are 11 missing values in TotalCharges column




```
In [22]: data.loc[data["TotalCharges"].isnull()== True]
```

Out[22]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Multi
488	4472-LVYGI	Female	0	Yes	Yes	0	No	^
753	3115-CZMZD	Male	0	No	Yes	0	Yes	
936	5709-LVOEQ	Female	0	Yes	Yes	0	Yes	
1082	4367-NUYAO	Male	0	Yes	Yes	0	Yes	
1340	1371-DWPAZ	Female	0	Yes	Yes	0	No	^
3331	7644-OMVMY	Male	0	Yes	Yes	0	Yes	
3826	3213-VVOLG	Male	0	Yes	Yes	0	Yes	
4380	2520-SGTTA	Female	0	Yes	Yes	0	Yes	
5218	2923-ARZLG	Male	0	Yes	Yes	0	Yes	
6670	4075-WKNIU	Female	0	Yes	Yes	0	Yes	
6754	2775-SEFEE	Male	0	No	Yes	0	Yes	

11 rows × 21 columns



## Missing value treatment

since the % of these records compared to total dataset is very low i.e. 0.15% it is safe to ignore them from further processing

```
In [23]: # Removing the missing values
data.dropna(how = 'any', inplace = True)
```

Dividing the customers into bins based on tenure e.g. for tenure < 12 months assign a tenure group 1-12 for tenure group between 1 to 2 years tenure group of 13-24 and so on

```
In [24]: ▶ # get the max tenure  
print(data["tenure"].max())
```

72

```
In [25]: ▶ # Group the tenure in the bins of 12 months  
labels = [f"{i}-{i+11}" for i in range(1,72,12)]
```

```
In [26]: ▶ labels
```

```
Out[26]: ['1-12', '13-24', '25-36', '37-48', '49-60', '61-72']
```

```
In [27]: ▶ data["tenure_group"] = pd.cut(data["tenure"], range(1,80,12), right = False)
```

```
In [28]: ▶ data["tenure_group"].value_counts()
```


```
Out[28]: 1-12      2175  
61-72      1407  
13-24      1024  
25-36       832  
49-60       832  
37-48       762  
Name: tenure_group, dtype: int64
```

Remove columns not required for processing

```
In [29]: ▶ # drop column customerID and tenure  
data.drop(columns= ["customerID", "tenure"], axis = 1, inplace = True)
```

## Data Exploration

```
In [30]: ▶ # Univariate analysis
for i, predictor in enumerate(data.drop(columns = ["Churn", "TotalCharges"]):
    plt.figure(i)
    sns.countplot(data = data, x = predictor, hue = "Churn")
```



```
In [31]: ▶ # Convert the target variable "Churn" in a binary numeric variable i.e. yes
```

```
In [32]: ▶ data["Churn"] = np.where(data.Churn == "Yes", 1, 0)
```


```
In [33]: ▶ data.head()
```

```
Out[33]:
```

	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService
0	Female	0	Yes	No	No	No phone service	DSL
1	Male	0	No	No	Yes	No	DSL
2	Male	0	No	No	Yes	No	DSL
3	Male	0	No	No	No	No phone service	DSL
4	Female	0	No	No	Yes	No	Fiber optic

```
In [34]: ▶ data.corr()["Churn"].sort_values(ascending= False).plot(kind = "bar")  
plt.show()
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_20092\2440828981.py:1: Future Warning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.  
data.corr()["Churn"].sort\_values(ascending= False).plot(kind = "bar")



```
In [35]: ▶ # converting all categorical variables into dummy variables
```

```
In [36]: ▶ data_dummy = pd.get_dummies(data)
```

```
In [37]: data_dummy.head(5)
```


Out[37]:

	SeniorCitizen	MonthlyCharges	TotalCharges	Churn	gender_Female	gender_Male	Partn
0	0	29.85	29.85	0	1	0	
1	0	56.95	1889.50	0	0	1	
2	0	53.85	108.15	1	0	1	
3	0	42.30	1840.75	0	0	1	
4	0	70.70	151.65	1	1	0	

5 rows × 51 columns

```
In [38]: # Relationship between monthly charges and total charges
sns.lmplot(data = data_dummy, x = "MonthlyCharges", y = "TotalCharges", fi
```

Out[38]: <seaborn.axisgrid.FacetGrid at 0x19547e38c90>




Total Charges increases as monthly charges increases as expected

Churn by Monthly Charges and Total Charges

```
In [39]: ▶ #Build a co relation of all predictors with churn
plt.figure(figsize= (20,8))
data_dummy.corr()["Churn"].sort_values(ascending = False).plot(kind = "bar")
```

Out[39]: <Axes: >




```
In [40]: ▶ # Derived insights
# High churn is seen in case of month to month contracts, no_online securi
# and fiber optics internet

# Low churn is seen in case of Long term contracts, subscriptions without
# 5+ years

# factors like gender, availability of phone service and # of multiple lin
```

```
In [41]: plt.figure(figsize=(12,12))
sns.heatmap(data_dummy.corr(), cmap = "Paired")
```

Out[41]: <Axes: >



## Bivariate Analysis


```
In [45]: new_data_target0 = data.loc[data["Churn"]== 0]
```

```
In [46]: new_data_target1 = data.loc[data["Churn"]== 1]
```

```
In [48]: ▶ def uniplot(df,col,title, hue = None):
sns.set_style("whitegrid")
sns.set_context("talk")
plt.rcParams["axes.labelsize"] = 20
plt.rcParams["axes.titlesize"] = 22
plt.rcParams["axes.titlepad"] = 30


temp = pd.Series(data = hue)
fig, ax = plt.subplots()
width = len(df[col].unique()) +7 + 4*len(temp.unique())
fig.set_size_inches(width,8)
plt.xticks(rotation = 45)
plt.yscale("log")
plt.title(title)
ax = sns.countplot(data = df, x = col, order = df[col].value_counts().
plt.show()
```

```
In [52]: ▶ uniplot(new_data_target1, col = "Partner", title = "Distribution of Gender
```







```
In [53]: ▶ uniplot(new_data_target1, col = "Contract", title = "Distribution of conta
```



```
In [54]: ▶ uniplot(new_data_target1, col = "PaymentMethod", title = "Distribution of
```



```
In [55]: for i, predictor in enumerate(data.drop(columns = ["Churn", "TotalCharges"]):
plt.figure(i)
sns.countplot(data = data, x = predictor, hue = "Churn")
```



2. Convert the target variable "Churn" in a binary numerical variable i.e. yes = 1 and no = 0

```
In [57]: data["Churn"] = np.where(data.Churn == "Yes", 1, 0)
```

```
In [58]: data.head()
```

```
Out[58]:
```

	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService
0	Female	0	Yes	No	No	No phone service	DSL
1	Male	0	No	No	Yes	No	DSL
2	Male	0	No	No	Yes	No	DSL
3	Male	0	No	No	No	No phone service	DSL
4	Female	0	No	No	Yes	No	Fiber optic

Conclusion These are the some of the quick insights from this exercise

1. Electronic check medium are the highest churners
2. contract type-monthly customers are more likely to churn because of no contract terms as they are free to go customers
3. Non senior citizens are high churners

In [65]: ▶

In [ ]: ▶